

Qwen3.5-9B English: every mitigation on the same 450 sets

Military status, implicit condition, 5 variants per set. Bars: 95% candidate-clustered bootstrap CI. Lexical scrubbing is an upper bound, not a deployable method.

