

lapa-v0.1.2-instruct EN: gender — every mitigation against the baseline on the same rows

Mitigation	Condition	Unstable sets % [95% CI]	Δ pp	Fixed : broken	p	MAD pp ↓	AR range pp ↓	Cohen's h ↓	Inconsist. % ↓	Hire %	Utility % ↑	Mean FS	Attr. mentioned %
baseline (no mitigation)	implicit	11.11 [7.78, 14.44]				1.41	8.00	0.11	1.98	84.13	48.04	0.56	0.43
SFT (LoRA, merged)	implicit	2.89 [1.56, 4.44]	-8.22	50 : 13	3e-06	0.22 (1.41)	0.89 (8.00)	0.01 (0.11)	0.50 (1.98)	17.72 (84.13)	80.70 (48.04)	0.63 (0.56)	0.27 (0.43)
baseline (no mitigation)	explicit	20.22 [15.78, 24.89]				2.12	17.33	0.10	2.73	82.16	50.18	0.57	0.36
SFT (LoRA, merged)	explicit	0.89 [0.22, 2.00]	-19.33	90 : 3	3e-23	0.11 (2.12)	0.67 (17.33)	0.00 (0.10)	0.19 (2.73)	17.54 (82.16)	80.37 (50.18)	0.63 (0.57)	0.31 (0.36)

Values: mitigated run, baseline in brackets, both on exactly the rows the run evaluated. Unstable sets: share with its own 95% CI (bootstrap over candidates); Δ pp: change against the baseline; bold coloured Δ = p < 0.05 (exact sign test on fixed vs broken sets, uncorrected). A run with parse failures is compared on its complete sets only; its own baseline on those sets is shown under its value. Italic ♦ = best value in the column within its condition (real mitigations only; hire rate and FS have no better direction).

Fixed : broken — each counterfactual set (one candidate–job pair, one condition, all attribute variants) paired with the same set at baseline. Fixed: unstable at baseline (the attribute alone changed the decision), stable after the mitigation. Broken: stable at baseline, unstable after — the mitigation introduced a new dependence on the attribute. Δ unstable sets = (broken – fixed) / sets; e.g. 58 : 57 means the method fixes about as many sets as it breaks.

* Lexical scrubbing is an oracle upper bound, not a comparable mitigation: it removes the attribute with the study's own injection templates, so every variant becomes character-for-character the attribute-free CV (100% of rows) and the model sees one identical prompt — zero instability by construction. Real CVs do not use our templates; read this row only as the ceiling. LLM scrubbing, which must find the attribute itself, is the realistic version.