

lapa-v0.1.2-instruct UK: gender — every mitigation against the baseline on the same rows

Mitigation	Condition	Unstable sets % [95% CI]	Δ pp	Fixed : broken	p	MAD pp ↓	AR range pp ↓	Cohen's h ↓	Inconsist. % ↓	Hire %	Utility % ↑	Mean FS	Attr. mentioned %
baseline (no mitigation)	implicit	13.62 [10.04, 17.19]				1.59	8.02	0.09	2.46	82.94	49.58	0.58	0.41
SFT (LoRA, merged)	implicit	3.35 [1.78, 5.13]	-10.27	61 : 15	1e-07	0.20 (1.59)	1.11 (8.02)	0.01 (0.09)	0.57 (2.46)	16.70 (82.94)	77.46 (49.58)	0.66 (0.58)	0.09 (0.41)
baseline (no mitigation)	explicit	30.04 [24.32, 35.73]				4.77	24.59	0.11	5.24	74.38	55.66	0.60	0.28
SFT (LoRA, merged)	explicit	2.69 [1.35, 4.26]	-27.35	134 : 12	3e-27	0.36 (4.77)	1.56 (24.59)	0.01 (0.11)	0.56 (5.24)	16.00 (74.38)	77.31 (55.66)	0.65 (0.60)	0.09 (0.28)

Values: mitigated run, baseline in brackets, both on exactly the rows the run evaluated. Unstable sets: share with its own 95% CI (bootstrap over candidates); Δ pp: change against the baseline; bold coloured Δ = p < 0.05 (exact sign test on fixed vs broken sets, uncorrected). A run with parse failures is compared on its complete sets only; its own baseline on those sets is shown under its value. Italic ♦ = best value in the column within its condition (real mitigations only; hire rate and FS have no better direction).

Fixed : broken — each counterfactual set (one candidate–job pair, one condition, all attribute variants) paired with the same set at baseline. Fixed: unstable at baseline (the attribute alone changed the decision), stable after the mitigation. Broken: stable at baseline, unstable after — the mitigation introduced a new dependence on the attribute. Δ unstable sets = (broken – fixed) / sets; e.g. 58 : 57 means the method fixes about as many sets as it breaks.

* Lexical scrubbing is an oracle upper bound, not a comparable mitigation: it removes the attribute with the study's own injection templates, so every variant becomes character-for-character the attribute-free CV (100% of rows) and the model sees one identical prompt — zero instability by construction. Real CVs do not use our templates; read this row only as the ceiling. LLM scrubbing, which must find the attribute itself, is the realistic version.