

lapa-v0.1.2-instruct EN: military status — every mitigation against the baseline on the same rows

Mitigation	Condition	Unstable sets % [95% CI]	Δ pp	Fixed : broken	p	MAD pp ↓	AR range pp ↓	Cohen's h ↓	Inconsist. % ↓	Hire %	Utility % ↑	Mean FS	Attr. mentioned %
baseline (no mitigation)	implicit	32.00 [26.67, 37.78]				8.75	24.89	0.20	8.93	75.82	53.96	0.56	6.71
SFT (LoRA, merged)	implicit	3.11 [1.78, 4.67]	-28.89	141 : 11	7e-30	0.32 (8.75)	0.89 (24.89)	0.01 (0.20)	0.84 (8.93)	19.60 (75.82)	80.49 (53.96)	0.63 (0.56)	0.18 (6.71)
baseline (no mitigation)	explicit	23.39 [18.89, 28.29]				6.04	19.76	0.15	6.22	82.40	49.31	0.55	2.49
SFT (LoRA, merged)	explicit	2.45 [1.11, 4.00]	-20.94	103 : 9	2e-21	0.46 (6.04)	1.56 (19.76)	0.01 (0.15)	0.67 (6.22)	18.53 (82.40)	81.02 (49.31)	0.64 (0.55)	0.13 (2.49)

Values: mitigated run, baseline in brackets, both on exactly the rows the run evaluated. Unstable sets: share with its own 95% CI (bootstrap over candidates); Δ pp: change against the baseline; bold coloured Δ = p < 0.05 (exact sign test on fixed vs broken sets, uncorrected). A run with parse failures is compared on its complete sets only; its own baseline on those sets is shown under its value. Italic ♦ = best value in the column within its condition (real mitigations only; hire rate and FS have no better direction).

Fixed : broken — each counterfactual set (one candidate–job pair, one condition, all attribute variants) paired with the same set at baseline. Fixed: unstable at baseline (the attribute alone changed the decision), stable after the mitigation. Broken: stable at baseline, unstable after — the mitigation introduced a new dependence on the attribute. Δ unstable sets = (broken – fixed) / sets; e.g. 58 : 57 means the method fixes about as many sets as it breaks.

* Lexical scrubbing is an oracle upper bound, not a comparable mitigation: it removes the attribute with the study's own injection templates, so every variant becomes character-for-character the attribute-free CV (100% of rows) and the model sees one identical prompt — zero instability by construction. Real CVs do not use our templates; read this row only as the ceiling. LLM scrubbing, which must find the attribute itself, is the realistic version.