

lapa-v0.1.2-instruct UK: military status — every mitigation against the baseline on the same rows

Mitigation	Condition	Unstable sets % [95% CI]	Δ pp	Fixed : broken	p	MAD pp ↓	AR range pp ↓	Cohen's h ↓	Inconsist. % ↓	Hire %	Utility % ↑	Mean FS	Attr. mentioned %
baseline (no mitigation)	implicit	25.39 [20.44, 30.51]				6.85	20.00	0.16	7.56	73.42	56.91	0.59	2.84
SFT (LoRA, merged)	implicit	2.90 [1.34, 4.67]	-22.49	113 : 12	9e-22	0.25 (6.85)	0.89 (20.00)	0.01 (0.16)	0.67 (7.56)	16.53 (73.42)	77.07 (56.91)	0.65 (0.59)	0.09 (2.84)
baseline (no mitigation)	explicit	17.63 [13.23, 22.54]				4.71	13.97	0.11	5.08	79.16	52.49	0.59	1.42
SFT (LoRA, merged)	explicit	1.34 [0.44, 2.46]	-16.29	79 : 6	2e-17	0.32 (4.71)	0.89 (13.97)	0.01 (0.11)	0.44 (5.08)	16.49 (79.16)	76.93 (52.49)	0.65 (0.59)	0.00 (1.42)

Values: mitigated run, baseline in brackets, both on exactly the rows the run evaluated. Unstable sets: share with its own 95% CI (bootstrap over candidates); Δ pp: change against the baseline; bold coloured Δ = p < 0.05 (exact sign test on fixed vs broken sets, uncorrected). A run with parse failures is compared on its complete sets only; its own baseline on those sets is shown under its value. Italic ♦ = best value in the column within its condition (real mitigations only; hire rate and FS have no better direction).

Fixed : broken — each counterfactual set (one candidate–job pair, one condition, all attribute variants) paired with the same set at baseline. Fixed: unstable at baseline (the attribute alone changed the decision), stable after the mitigation. Broken: stable at baseline, unstable after — the mitigation introduced a new dependence on the attribute. Δ unstable sets = (broken – fixed) / sets; e.g. 58 : 57 means the method fixes about as many sets as it breaks.

* Lexical scrubbing is an oracle upper bound, not a comparable mitigation: it removes the attribute with the study's own injection templates, so every variant becomes character-for-character the attribute-free CV (100% of rows) and the model sees one identical prompt — zero instability by construction. Real CVs do not use our templates; read this row only as the ceiling. LLM scrubbing, which must find the attribute itself, is the realistic version.