

lapa-v0.1.2-instruct UK: religion — every mitigation against the baseline on the same rows

Mitigation	Condition	Unstable sets % [95% CI]	Δ pp	Fixed : broken	p	MAD pp ↓	AR range pp ↓	Cohen's h ↓	Inconsist. % ↓	Hire %	Utility % ↑	Mean FS	Attr. mentioned %
baseline (no mitigation)	implicit	20.67 [16.22, 24.94]				3.84	14.96	0.11	4.74	68.84	59.85	0.61	1.28
SFT (LoRA, merged)	implicit	3.37 [1.80, 5.15]	-17.30	92 : 15	1e-14	0.68 (3.84)	2.00 (14.96)	0.02 (0.11)	0.99 (4.74)	16.77 (68.84)	77.11 (59.85)	0.65 (0.61)	0.00 (1.28)
baseline (no mitigation)	explicit	24.05 [19.51, 29.02]				5.32	21.21	0.22	5.77	65.19	61.70	0.62	0.12
SFT (LoRA, merged)	explicit	1.11 [0.22, 2.22]	-22.94	108 : 5	3e-26	0.20 (5.32)	0.67 (21.21)	0.01 (0.22)	0.22 (5.77)	15.95 (65.19)	77.09 (61.70)	0.65 (0.62)	0.00 (0.12)

Values: mitigated run, baseline in brackets, both on exactly the rows the run evaluated. Unstable sets: share with its own 95% CI (bootstrap over candidates); Δ pp: change against the baseline; bold coloured Δ = p < 0.05 (exact sign test on fixed vs broken sets, uncorrected). A run with parse failures is compared on its complete sets only; its own baseline on those sets is shown under its value. Italic ♦ = best value in the column within its condition (real mitigations only; hire rate and FS have no better direction).

Fixed : broken — each counterfactual set (one candidate–job pair, one condition, all attribute variants) paired with the same set at baseline. Fixed: unstable at baseline (the attribute alone changed the decision), stable after the mitigation. Broken: stable at baseline, unstable after — the mitigation introduced a new dependence on the attribute. Δ unstable sets = (broken – fixed) / sets; e.g. 58 : 57 means the method fixes about as many sets as it breaks.

* Lexical scrubbing is an oracle upper bound, not a comparable mitigation: it removes the attribute with the study's own injection templates, so every variant becomes character-for-character the attribute-free CV (100% of rows) and the model sees one identical prompt — zero instability by construction. Real CVs do not use our templates; read this row only as the ceiling. LLM scrubbing, which must find the attribute itself, is the realistic version.